
Proximal Policy Optimization (PPO)

Shahzeb Aamir

Content:

1. Motivation for PPO
2. PPO Goals
3. PPO Algorithm in detail
4. Some cool implementations

Motivation

Training Data is not “static”

Training data is dependent upon current policy

High Sensitivity to Hyper-Parameters

If we get one “bad” policy, one outlier, it can end up in a fire pit. Or a very slow progress

PPO Goals

Developed by folks at OpenAI, based on Policy Gradient Methods

John Schulman et al., Proximal Policy Optimization Algorithms (2017 Paper)

1. Ease of implementation (Easy Code)
2. *Sample Efficiency ("Online Policy Method")*
3. *Robust (Less sensitive to Hyper Parameters)*

Policy Gradient Loss

1. Vanilla Policy Gradient

$$L^{PG}(\theta) = \hat{\mathbb{E}}_t \left[\log \pi_{\theta}(a_t | s_t) \hat{A}_t \right].$$

1. Discounted Rewards (Return)
2. Baseline Estimate (Value Function)

Advantage: Relative value of the action

Policy: A neural network that takes in observed state and gives probability of actions

Trusted Region

New policy and old policy is not too different

This is the basis of PPO

$$\begin{aligned} & \underset{\theta}{\text{maximize}} && \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)} \hat{A}_t \right] \\ & \text{subject to} && \hat{\mathbb{E}}_t [\text{KL}[\pi_{\theta_{\text{old}}}(\cdot | s_t), \pi_{\theta}(\cdot | s_t)]] \leq \delta. \end{aligned}$$

Value Function: Guess the final return starting from the current state

Usually a neural network: A noisy estimate

PPO

1. Crux of PPO

The surrogate policy loss function:

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}$$

This probability ratio > 1 : If the action is more likely now than it was before gradient update

0 $<$ This probability ratio < 1 : If the action is less likely now than it was before gradient update

PPO

1. If advantage is positive then the probability of taking that action increases but not too much (that it clips)
2. If advantage is negative then the probability of taking that action decreases but not too much




Figure 1: Plots showing one term (i.e., a single timestep) of the surrogate function L^{CLIP} as a function of the probability ratio r , for positive advantages (left) and negative advantages (right). The red circle on each plot shows the starting point for the optimization, i.e., $r = 1$. Note that L^{CLIP} sums many of these terms.

PPO

1. Final Objective function:

$$L_t^{CLIP+VF+S}(\theta) = \hat{\mathbb{E}}_t [L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t)]$$

Algorithm

Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1,2,... do  
  for actor=1,2,...,N do  
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps  
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$   
  end for  
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$   
   $\theta_{\text{old}} \leftarrow \theta$   
end for
```

PPO

Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1,2,... do
  for actor=1,2,...,N do
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$ 
  end for
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$ 
   $\theta_{\text{old}} \leftarrow \theta$ 
end for
```

$$\hat{A}_t = \delta_t + (\gamma\lambda)\delta_{t+1} + \dots + \dots + (\gamma\lambda)^{T-t+1}\delta_{T-1},$$

where $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

PPO implementation

<https://openai.com/blog/openai-baselines-ppo/>

<https://github.com/higgsfield/RL-Adventure-2/blob/master/3.ppo.ipynb>

<https://www.youtube.com/watch?v=WxQfQW48A4A&list=PLB79uOaPEEU6uU1-Pfaqr08RTTzhyB8hu&index=3>



Skolkovo Institute of Science and Technology

Thank You!